

Multinomial Logit Model

Your Name

2025-05-19

This assignment explores two methods for estimating the MNL model: (1) via Maximum Likelihood, and (2) via a Bayesian approach using a Metropolis-Hastings MCMC algorithm.

1. Likelihood for the Multi-nomial Logit (MNL) Model

Suppose we have $i = 1, \dots, n$ consumers who each select exactly one product j from a set of J products. The outcome variable is the identity of the product chosen $y_i \in \{1, \dots, J\}$ or equivalently a vector of $J - 1$ zeros and 1 one, where the 1 indicates the selected product. For example, if the third product was chosen out of 3 products, then either $y = 3$ or $y = (0, 0, 1)$ depending on how we want to represent it. Suppose also that we have a vector of data on each product x_j (eg, brand, price, etc.).

We model the consumer's decision as the selection of the product that provides the most utility, and we'll specify the utility function as a linear function of the product characteristics:

$$U_{ij} = x'_j \beta + \epsilon_{ij}$$

where ϵ_{ij} is an i.i.d. extreme value error term.

The choice of the i.i.d. extreme value error term leads to a closed-form expression for the probability that consumer i chooses product j :

$$\mathbb{P}_i(j) = \frac{e^{x'_j \beta}}{\sum_{k=1}^J e^{x'_k \beta}}$$

For example, if there are 3 products, the probability that consumer i chooses product 3 is:

$$\mathbb{P}_i(3) = \frac{e^{x'_3 \beta}}{e^{x'_1 \beta} + e^{x'_2 \beta} + e^{x'_3 \beta}}$$

A clever way to write the individual likelihood function for consumer i is the product of the J probabilities, each raised to the power of an indicator variable (δ_{ij}) that indicates the chosen product:

$$L_i(\beta) = \prod_{j=1}^J \mathbb{P}_i(j)^{\delta_{ij}} = \mathbb{P}_i(1)^{\delta_{i1}} \times \dots \times \mathbb{P}_i(J)^{\delta_{iJ}}$$

Notice that if the consumer selected product $j = 3$, then $\delta_{i3} = 1$ while $\delta_{i1} = \delta_{i2} = 0$ and the likelihood is:

$$L_i(\beta) = \mathbb{P}_i(1)^0 \times \mathbb{P}_i(2)^0 \times \mathbb{P}_i(3)^1 = \mathbb{P}_i(3) = \frac{e^{x'_3\beta}}{\sum_{k=1}^3 e^{x'_k\beta}}$$

The joint likelihood (across all consumers) is the product of the n individual likelihoods:

$$L_n(\beta) = \prod_{i=1}^n L_i(\beta) = \prod_{i=1}^n \prod_{j=1}^J \mathbb{P}_i(j)^{\delta_{ij}}$$

And the joint log-likelihood function is:

$$\ell_n(\beta) = \sum_{i=1}^n \sum_{j=1}^J \delta_{ij} \log(\mathbb{P}_i(j))$$

2. Simulate Conjoint Data

We will simulate data from a conjoint experiment about video content streaming services. We elect to simulate 100 respondents, each completing 10 choice tasks, where they choose from three alternatives per task. For simplicity, there is not a “no choice” option; each simulated respondent must select one of the 3 alternatives.

Each alternative is a hypothetical streaming offer consistent of three attributes: (1) brand is either Netflix, Amazon Prime, or Hulu; (2) ads can either be part of the experience, or it can be ad-free, and (3) price per month ranges from \$4 to \$32 in increments of \$4.

The part-worths (ie, preference weights or beta parameters) for the attribute levels will be 1.0 for Netflix, 0.5 for Amazon Prime (with 0 for Hulu as the reference brand); -0.8 for included advertisements (0 for ad-free); and $-0.1 \times \text{price}$ so that utility to consumer i for hypothetical streaming service j is

$$u_{ij} = (1 \times \text{Netflix}_j) + (0.5 \times \text{Prime}_j) + (-0.8 \times \text{Ads}_j) - 0.1 \times \text{Price}_j + \varepsilon_{ij}$$

where the variables are binary indicators and ε is Type 1 Extreme Value (ie, Gumble) distributed.

The following code provides the simulation of the conjoint data.

i Note

```

# set seed for reproducibility
# set.seed(123)

# # define attributes
# brand <- c("N", "P", "H") # Netflix, Prime, Hulu
# ad <- c("Yes", "No")
# price <- seq(8, 32, by=4)

# # generate all possible profiles
# profiles <- expand.grid(
#   brand = brand,
#   ad = ad,
#   price = price
# )
# m <- nrow(profiles)

# # assign part-worth utilities (true parameters)
# b_util <- c(N = 1.0, P = 0.5, H = 0)
# a_util <- c(Yes = -0.8, No = 0.0)
# p_util <- function(p) -0.1 * p

# # number of respondents, choice tasks, and alternatives per task
# n_peeps <- 100
# n_tasks <- 10
# n_alts <- 3

# # function to simulate one respondent's data
# sim_one <- function(id) {

#   datlist <- list()

#   # loop over choice tasks
#   for (t in 1:n_tasks) {

#     # randomly sample 3 alts (better practice would be to use a design)
#     dat <- cbind(resp=id, task=t, profiles[sample(m, size=n_alts), ])

#     # compute deterministic portion of utility
#     dat$v <- b_util[dat$brand] + a_util[dat$ad] + p_util(dat$price) |> round(10)

#     # add Gumbel noise (Type I extreme value)
#     dat$e <- -log(-log(runif(n_alts)))
#     dat$u <- dat$v + dat$e

#     # identify chosen alternative
#     dat$choice <- as.integer(dat$u == max(dat$u))

#     # store task
#     datlist[[t]] <- dat
#   }

#   # combine all tasks for one respondent
#   do.call(rbind, datlist)

```

3. Preparing the Data for Estimation

4. Estimation via Maximum Likelihood

5. Estimation via Bayesian Methods

6. Discussion